

Reasoning about Concurrency in (RT) Java

A description of past work on reasoning about
concurrent Java and ongoing in the
EU FP6 Artimis “CHARTER” project

Joe Kiniry

Outline

- Java vs. JavaCard vs. RT Java
- reasoning about concurrency in the past
- modern models of concurrency
- what do developers (do/understand)?
- recent work and next steps

Concurrency in Java

Early Java

- only two classes: *Thread* and *ThreadGroup*
- *ThreadGroups* are sets of threads, arranged hierarchically, with a pseudo-multicast façade
- *synchronized* and *volatile* keywords exist, but the latter is ignored
- hand-waving about scheduling: “Every thread has a priority. Threads with higher priority are executed in preference to threads with lower priority.”
- thread has several surprising methods: *checkAccess*, *destroy*, *interrupt*, *setPriority*, *stop*

Safety

- every object contains a monitor
- a thread attempts to lock an object's monitor either via a call to a synchronized method or an explicit synchronized block
- recall that classes are represented by singleton objects of type *Class*, thus synchronized static methods lock these singletons

Safety Problems

- no guarantees on lock ordering, fairness, complex semantics on lock release, etc.
- no support for identifying race conditions
- loose and very complex semantics on memory consistency (Manson, Pugh, and Adve in POPL'05)
- reordering during and after JIT permitted

Liveness

- threads communicate via calls to low-level methods like *sleep*, *wait*, and *notify/notifyAll*
- no support to identifying or avoiding deadlock or livelock
- no semantics for priorities and scheduling

Communication

- superposition via volatile shared variables with non-atomic updates for nearly all primitive types
- wait/notify patterns often misused
- manual encoding of barriers, callbacks, etc.

Teenager Java (1.2–1.4)

- *Thread* now has access to its *ClassLoader* and gets the *holdLock* method
- *ThreadGroup* now has *interrupt* method
- *ThreadLocal* class introduced
- *Thread*'s *suspend*, *resume*, and *stop* methods are all deprecated
- Doug Lea starts work on his Java concurrency library (*EDU.oswego.cs.dl.util.concurrent*)

Mature Java (1.5)

- *Thread* gets structured access to stack traces, ID, state, and exception handlers
- thread state is now explicit and exactly one of the following states:
 - *new, runnable, blocked, waiting, timed_waiting, terminated*
- *ThreadLocal* gets a *remove* method
- Doug's concurrency library standardized

The Java Concurrency API

- lifts the level of abstraction away from raw threads, monitors, and synchronized regions
- new constructs available include:
 - atomic wrapper classes for some primitive types and references, barriers, concurrent collections, conditions, copy-on-write collections, executors, futures, latches, (pairs of) locks, queues, semaphores, threadpools, etc.
- implementation is VM-invariant and somewhat lock-free

Java Variants

- JavaCard
 - no concurrency or floating point and a different memory model than normal Java (no allocation, permanent and transient memory)
- MIDP
 - normal, concurrent Java, but on small devices
- Real-time Java (RT Java)
 - soft & hard real-time with priorities with deterministic scheduling and mixed thread model
 - *Thread's* API becomes well-defined again
 - triplet (scoped/immortal/heap) memory model

Reasoning about Concurrency

Early Efforts

- Model Checking
 - Bandera and Bogor from Corbett, Dwyer, Hatcliff, Laubach, Pasareanu, Robby, and Zheng (ICSE'00—CAV'05)
 - Java PathFinder from Havelund and Pressburger (SPIN'99—TACAS'07)
- Proof Systems
 - Ábrahám, de Boer, de Roever and Steffen (CONCUR'00—Fund. Info.'08)

Finding Key Abstractions

- Race Condition Checking from Flanagan and Freund (ESOP'99—PLDI'00)
- Atomicity from Flanagan, Freund, and Qadeer (PLDI'03—PLDI'08)
- Immutability (Ernst and many colleagues OOPSLA'05—ESEC-FSE'07; Haack, Poll, Schäfer, and Schubert, ESOP'07; et al.)

“Purported” Tools

- there is a terrible dearth of released, supported tools in this research area
- dozens of tools for reasoning about concurrency have been discussed in papers, but have not been released to the research community
- IMO, this is not kosher

“Real” Tools

- JProbe’s deadlock analysis and race condition detection
- lightweight static checkers like PMD and FindBugs do conservative race analysis and concurrency anti-pattern detection
- RCC does type-based race condition checking
- ESC/Java2 does deadlock and locking discipline analysis

Recent Efforts

- the Mobius approach
 - ensure programs are *properly synchronized* through the use of ESC/Java2 and RCC, then reason about program sequentially
- deductive verification from Beckert and Klebenov (threading) in the KeY group
- the separation logic camp
 - Berdine, Bierman, Calcagno, Distefano, Huisman, Hurlin, Jacobs, O'Hearn, Parkinson, Smans, Reynolds, Vafeiadis, Yang

Properly Synchronized

- several definitions of *properly synchronized* have been floated in the literature
- within Mobius, it meant:
 - a program has no race conditions
 - every shared variable is monitored by one or more locks
 - lock ordering is consistent
 - the locking discipline is respected

Specification

Constructs and Use

- JML permits one to specify through a small set of primitive constructs your own locking discipline
- if ESC/Java2 reports that your annotated code has no race conditions, then it *likely* has none
- if ESC/Java2 reports that your annotated code has no deadlocks, then it *likely* has none
- if RCC reports your annotated code has no race conditions, then it *definitely* has none

Locking Disciplines

- a locking discipline is *a means by which concurrently accessed data is guarded*
- a locking discipline answers the questions
 - what data is (not) encapsulated?
 - how does one access said data?
 - which constructs are used for access control and which ones for data?

Discipline Flavors

- conservative/pessimistic
- liberal/optimistic
- strict separation of data and access control
- the data is the access control
- hierarchical structuring/ownership
- permission granting

Expressing and Reasoning about Locking Disciplines

monitored

```
class C {  
    // this annotation means that read or write  
    // accesses to field 'f' must only happen  
    // after the containing object is locked  
    // Only legal for instance fields.  
    float f; //@ monitored  
  
    synchronized void m() { f = 1.0; }           // ok  
    synchronized static void n(C c)              // error!  
        { c.f = 1.0; }  
}  
  
// client code  
C c = new C();  
float g = c.f;                                   // error!  
c.f = 0.0;                                       // error!  
synchronized (c) { c.f = 1.0; } // ok
```


static monitors_for

```
class C {  
    // 'aLock' must be locked for all access to field 'i'  
    // If 'i' is static, then its lock must be static.  
    //@ monitors_for i <- aLock;  
    static int i;  
    static Object aLock;  
  
    static void m() {  
        synchronized(aLock) { i = 1; }    // ok  
    }  
    static void n() { i = 2; }              // error!  
}  
  
...  
C c = new C();  
C.i = 1.0;                                // error!  
synchronized(C.class) { c.d = 1.0; }      // error!  
synchronized(C.aLock) { c.d = 1.0; }      // ok
```

simple dynamic monitors_for

```
class C {
  double d;
  //@ monitors_for d <- this;
  // equivalent to just annotating with "monitored"

  synchronized void m() { d = 1.0; }           // ok
  void n() { d = 1.0; }                         // error!
  void o() { synchronized(this) { d = 1.0; } } // ok
  void p(D d) { synchronized(d.aLock) {
    d = 1.0; } } // error unless we can prove d.aLock
                  // always is equal to this.
  void q() { Object o = this;
    synchronized(o) { d = 1.0; } }             // ok
}

...
C c = new C();
c.d = 1.0;                                     // error!
synchronized(c) { c.d = 1.0; }                 // ok
```

richer dynamic monitors_for

```
class C {
    String aLock;
    double d;
    //@ monitors_for d <- aLock;

    synchronized void m() { d = 1.0; }           // error!
    void n() { d = 1.0; }                         // error!
    void o() { synchronized(this) { d = 1.0; } }  // error!
    void p() { synchronized(aLock) { d = 1.0; } } // ok
    void q() { Object o = anotherLock;
        synchronized(o) { d = 1.0; } }           // ok
}

...
C c = new C();
c.d = 1.0;                                       // error!
synchronized(c.aLock) { c.d = 1.0; }          // ok
```

Multiple Locks

```
class C {
    static Object aLock;
    String anotherLock;
    double d;
    //@ monitors_for d <- aLock, anotherLock;
    //@ axiom aLock < anotherLock;

    synchronized void m() { d = 1.0; }      // error!
    static void n() {
        synchronized(aLock) { d = 1.0; } }  // error!
    void o() { synchronized(anotherLock) {
        synchronized(aLock) { d = 1.0; } } } // error!
    }
    ...
    C c = new C();
    float g = c.d;                          // error, even for reads!
    synchronized (C.aLock) {
        synchronized (c.anotherLock) {
            c.d = 1.0; } }                  // ok
}
```

Subtleties

- locks are arbitrary objects
- must deal with aliasing of locks
- if the lock reference is null, one cannot lock
- field hiding matters
- spec-accessibility matters
- multiple monitors_for are permitted

Locksets

- $\backslash lockset$ is of type $\backslash LockSet$
- denotes the set of locks held by the current thread
- membership in locksets
 - expression of the form $S[L]$ where S is a spec-expr of type $\backslash LockSet$ and L is a spec-expr of ref-type denotes that L is a member of S
- $\backslash max(S)$ denotes the maximum element of S

Complex Example

```
public class Tree {
    public /*@ monitored */ Tree left, right;
    public /*@ monitored */ Thing contents;
    //@ axiom (\forall Tree t; t.left != null ==> t < t.left);
    //@ axiom (\forall Tree t; t.right != null ==> t < t.right);

    Tree(Thing c) {
        contents = c;
    }

    //@ requires \max(\lockset) <= this;
    public synchronized void visit() {
        contents.mungle();
        if (left != null) left.visit();
        if (right != null) right.visit();
    }
}
```

RCC Annotations

- variant of `monitors_for`, `guarded_by`, is used, e.g.,
`Type VarName /*# guarded_by */ LockSet`
- lockset names are used as shorthand, e.g.,
`/*# requires LockSet */`
`ReturnType MethodName(Args) {Body}`
- class-level annotations of locksets possible
`ClassName /*# {ghost Object LockSet} */ {ClassBody}`
`ClassName/*# {LockSet} */ VarName`
- thread-local and thread-shared are introduced
`/*# thread_local */ ClassDeclaration`
`/*# thread_shared */ ClassDeclaration`

**But what do developers
do and understand?**

Old Concurrent Java

- deadlock and system failure was pervasive before *Thread*'s methods were deprecated
- Flanagan et al. found that most Java methods are written *as if* they are atomic
- unsafe idioms and assumptions were and are pervasive (assumptions about immutability, atomicity, double-checking locks, lazy instantiation)

Open Questions

- Which concurrency constructs are used most often?
- What are the most common concurrency idioms witnessed in RT Java code?
- Does the Java concurrency library work at all in an RT Java setting?
- If it does not, why not, and how might it be changed to accommodate priorities?

Case Study

- a tool called the Histogram System has been formally specified (from formal requirements refined down to JML-annotated Java) to analyze developers' use of concurrency
- our plan is to...
 - analyze $> 100\text{M}$ NCSS of off-the-shelf Java
 - analyze all known public examples of RT Java (probably 10,000s of NCSS)

Recent Work and Next Steps

Reasoning about RT Java

- Java PathFinder has been applied to RT Java by Lindstrom, Mehlitz, and Visser (ATVA'05)
- timing and dataflow analysis for WCET and schedulability by many researchers
- deductive verification about memory from Engel in the KeY group
- *no work on deductive verification of safety and liveness, especially at the model level*

Recent Work

- concurrency annotations at the model level
- traceable formal refinement to JML
- foundation for reasoning about concurrent architectures at the model level

The Concurrency Semantic Property

- six kinds of annotations
- annotations permitted at class or method level
- class-level annotations are not inherited

<i>concurrent</i>	<i>sequential</i>	<i>guarded</i>
<i>failure</i>	<i>atomic</i>	<i>special</i>

Preliminaries

- *Feature Thread Count*: The thread count for a feature f of object o is the number of threads simultaneously executing f on object o .
- *Broken Object*: When an object is broken, the object's invariants and feature postconditions are no longer guaranteed.

Semantics via Examples: Concurrent and Sequential

```
interface I {  
    // @concurrent 5  
    void m();  
    // indicates that, if more than 5 threads enter m,  
    // then this is broken  
  
    // @sequential  
    void n();  
    // equivalent to @concurrent 1  
    ...  
}
```

Locks

```
...  
// @locks a, b, c  
void o();  
// o will attempt to acquire no more than  
// locks a, b, and c  
  
// @locks Void  
void p();  
// p will not attempt to acquire any locks  
  
// @locks *  
void q();  
// q may acquire any set of locks  
...
```

Guarded

```
...  
// @guarded a, b, c  
void r();  
// caller must hold a, b, and c prior to r being  
// executed, but need not acquire them before  
// calling r, and a, b, and c are released prior  
// to r terminating  
  
// @guarded 3  
void s();  
// equivalent to guarded Semaphore(3)  
...
```

Failure, Atomic, and Special

```
...  
// @concurrent 2  
// @failure MyException  
void t();  
// t throws an exception of type MyException  
// immediately when the number of threads  
// executing t is 2 and a new thread calls t  
  
// @atomic  
void t();  
// indicates that t is serializable wrt some  
// definition of atomicity  
  
// @special This method locks the database DB.  
void u();  
...
```

Specification Expressions

- *\broken*
- *\semaphore* and *\semaphore(descriptor)*
descriptor is a string representation of a method signature
- *\thread_count*
- *\thread_limit* and *\thread_limit(descriptor)*

Ex. Semantic Mapping

```
class C {  
    //@ invariant I;  
  
    //@ requires P; assignable A; ensures Q;  
    // @concurrency concurrent 4  
    public void m() {}  
}
```

```
class C {  
    //@ invariant !\broken ==> I;  
  
    /*@ normal_behavior  
        @ requires P; assignable A; ensures !\broken ==> Q;  
        @ also normal_behavior  
        @ requires \thread_count == \thread_limit;  
        @ assignable A;  
        @ ensures \broken; */  
    public void m() {}  
  
    /*@ invariant \typeof(this) == \type(C) ==>  
        @ \thread_limit("m()") == 4; */  
    //@ invariant 4 <= \thread_limit("m()");  
}
```

Sleeping Barber

```
class BARBER
  feature
    cut_hair
      -> c: CUSTOMER
        concurrency guarded Current
        ensure not c.needs_hair_cut
        ensure delta c.needs_hair_cut
    end
end
```



```

class BARBER_SHOP
  feature
    barber: BARBER
    num_seats: INTEGER
    ensure 0 < Result
  end
  get_hair_cut
    -> c: CUSTOMER
    concurrency concurrent (num_seats + 1)
    concurrency failure NO_SEATS
    concurrency locks barber
    ensure not c.needs_hair_cut; delta c.needs_hair_cut
    -- behavior: barber.cut_hair(c)
  end
  make
    -> the_barber: BARBER
    -> the_num_seats: INTEGER
    ensure barber = the_barber; num_seats = the_num_seats
  end
end

```

```

class CUSTOMER
  concurrency guarded Current
  feature
    shop: BARBER_SHOP
    ensure Result /= Void
  end
  needs_hair_cut: BOOLEAN
  set_hair_cut
    ensure not needs_hair_cut; delta needs_hair_cut
  end
  regular_activities
    require not needs_hair_cut
    ensure needs_hair_cut; delta needs_hair_cut
    -- behavior: whatever the customer does between haircuts
  end
  run
    concurrency locks shop.barber
    ensure delta needs_hair_cut
    -- behavior: repeatedly execute the sequence
    -- "regular_activities; retry shop.get_hair_cut until
    -- not needs_hair_cut"
  end
  make
    -> the_shop: BARBER_SHOP
    ensure shop = the_shop; not needs_hair_cut
  end
end
end

```

```
class MAIN
  feature
    make
      -> the_seats: INTEGER
      -> the_customers: INTEGER
      require 0 < the_seats; 0 < the_customers
      -- behavior: create a barber, create a barber shop with
      -- the_seats seats, create the_customers customers,
      -- and "run" each customer in a separate thread
    end
  end
end

class NO_SEATS inherit EXCEPTION end
```

Contributions and Limitations

- contributions
 - model-level concurrency annotations
 - design matches the typical set of concurrency patterns witnessed in code
- limitations
 - no notion of fairness is specified
 - unable to reason without a well-specified locking discipline
 - tool support still underway

Next Steps

- How to reason about concurrency at the architecture level?
- Which constructs do developers use (correctly) and how might they be modeled in our reasoning engines (HOL or SMT)?
- How might one introduce separation into JML? Should it be explicit or implicit?
- How might one introduce orchestration into JML? Does such help with respect to concurrency?